Houdini: Security Benchmarking of Container Confinement

Huzaifa Patel
Carleton University

David Barrera
Carleton University

Anil Somayaji
Carleton University

Abstract

While container-based workloads are now a standard part of our cloud infrastructure, container security remains a challenging problem. Container confinement is a particularly pressing problem, as without it, a single vulnerable application can be used to compromise entire clusters of containers. While we have many technologies that can be used to secure containers, currently there is no easy way to determine whether a given Docker configuration provides even a basic level of protection. Here we present `Houdini`, a security test suite for container confinement. Much as network scanning tools can help catch misconfigured firewalls, `Houdini` can check whether a given Docker container configuration properly enforces confinement. While it can be used to test deployable containers, `Houdini` is optimised for testing and comparing container confinement technologies. Here we present the motivation, design, implementation, and initial results on running `Houdini` on a set of different container configurations. By providing a benchmark framework by which container confinement technologies can be evaluated, we believe `Houdini` can help foster the development of next-generation container confinement technologies.

1 Introduction

Container-based workloads have become a cornerstone of modern cloud infrastructure, revolutionizing how applications are developed, deployed, and scaled, particularly within microservices architectures and DevOps workflows. Unlike hypervisor-based virtual machines (VMs), which each require a full operating system and rely on a hypervisor for hardware virtualization, containers are lightweight, share the host system's kernel, and incur minimal overhead. This efficiency makes them exceptionally well-suited for continuous integration and rapid scaling, allowing organizations to optimize resource utilization and reduce costs while ensuring quick, consistent deployments. Furthermore, containers enable seamless portability by packaging all userspace dependencies into a unified image, thereby avoiding the common problem of environment-specific discrepancies. Yet, despite containers' numerous benefits, they also introduce new security challenges, especially where container confinement is concerned.

Linux kernel features such as cgroups, namespaces, and mount points provide resource, process, and network isolation for containers. However, this separation is not foolproof. Kernel vulnerabilities, such as buffer overflows or privilege escalation flaws can allow a process to circumvent the very mechanisms designed to confine it. Misconfigurations can similarly open holes in container security, for instance by granting unnecessary capabilities, mounting sensitive host directories, or running containers in privileged mode. Finally, sophisticated attacks like container escapes exploit these weaknesses to break out of the container's resource and namespace isolation. Once an attacker gains elevated privileges or direct access to the host system, they can disrupt other containers, exfiltrate data, or compromise the entire infrastructure. Technologies like SELinux, AppArmor, and seccomp provide additional layers of security to avoid these issues, but are notoriously complex to configure and prone to misconfiguration. A best practices guide for container security [1] can help with configuring a container, but it offers no guarantee that the given configuration will perform as intended by the end-user. As container adoption skyrocketed across industries, particularly in multi-tenant cloud environments, these security challenges have become more pronounced.

To address the aforementioned challenges, the cloud industry has developed technologies like gVisor and Kata Containers [2], which isolate containers further by running them within lightweight VMs. While this approach is effective, it unnecessarily increases complexity and overhead. By developing methods to audit and validate confinement mechanisms, container security could be simplified by reducing reliance on such solutions.

Test suites play a crucial role in verifying that systems meet specified requirements and function correctly under defined conditions. They are widely used across computing disciplines, from performance benchmarking to security as-

sessment. For instance, the Standard Performance Evaluation Corporation (SPEC) [3] is best known for developing industry-standard benchmarks that evaluate CPU performance and power efficiency, while the Transaction Processing Performance Council (TPC) [4] establishes benchmarking standards for database performance. The Phoronix Test Suite [5] is a widely used benchmarking tool, particularly for Linux-based systems, providing automated performance evaluations across various hardware and software configurations. Other notable benchmarks include Geekbench, a cross-platform tool for assessing CPU and GPU performance, and IOzone, which specializes in file system I/O analysis.

Beyond performance testing, test suites are essential in software development to prevent regressions. Many organizations mandate that all code changes pass a comprehensive suite of tests before integration. In production environments, test suites help maintain reliability through uptime monitoring, ensuring that services remain operational. Additionally, specialized security-focused test suites play a vital role in detecting vulnerabilities in deployed systems. Tools such as network scanners aid in identifying misconfigurations, outdated software versions, and unintended service exposures, helping organizations proactively mitigate security risks.

Despite the widespread use of test suites in the aforementioned areas, a significant gap remains in the context of containers. There are currently no standardized or widely adopted test suites designed specifically to evaluate container confinement, functionality, or isolation effectiveness. This absence leaves organizations to rely on ad hoc methods, which are often inadequate for uncovering subtle misconfigurations or ensuring robust confinement. Given the growing reliance on containerized environments in modern infrastructure, the development of comprehensive test suites tailored for containers is imperative. Such test suites would not only address these critical gaps but also provide a reliable foundation for improving trust and security in containerized systems.

In this paper, we describe `Houdini`, the first test suite for verifying container confinement. Given a Docker container configuration (including both the kernel version and the docker version), `Houdini` will instantiate the container in a standalone QEMU-based VM and perform multiple tests (or tricks, as we call them) to evaluate whether the container is susceptible to various forms of misconfigurations or failures. These tests explore vulnerabilities that could disrupt container functionality, such as, but not limited to improper resource allocation, inadequate isolation, or insecure filesystem setups. Houdini is written in Python and is easily extensible, providing an extension language for writing tricks. `Houdini` is designed to assess and enhance the security of containerized environments. It tests the isolation of containers, ensuring they remain properly confined and separated from the host system and other containers. Houdini specializes in identifying container-specific vulnerabilities, misconfigurations, and ensuring adherence to security best practices.

A standard methodology for assessing container confinement is crucial due to the complexity of modern systems and the intricate vulnerabilities they present. A formal approach involving rigorously evaluating and analyzing the security of the container system might seem ideal, however, it is impractical because container environments are highly complex, with interactions between kernels, runtimes, orchestration tools, and external dependencies that are difficult to model comprehensively. Moreover, many vulnerabilities stem from subtle details that a formal model might overlook. Instead, an empirical approach offers a practical alternative. Though it will never be exhaustive, the goal is not comprehensive coverage but rather establishing a baseline set of tests that demonstrate effective container confinement. This approach is achievable and can evolve over time as new vulnerabilities are identified, allowing the methodology to remain current. Starting with tests for isolation, privilege escalation prevention, and resource limitation enforcement provides a solid foundation. By focusing on demonstrating security effectiveness in critical areas, this methodology can enhance confidence in container confinement while enabling continuous improvement and adaptation to the evolving threat landscape.

In this paper, we describe the motivation, design, and implementation of `Houdini`. We also present the results of case studies by showing how Houdini can be used to detect misconfigured containers. Our hope with Houdini is that it will help with the deployment of better confined containers and will support the development of new technologies to more reliably confine containers.

The rest of this paper proceeds as follows. In Section 2, we explain the container confinement problem and associated technologies. Section 3 describes related work. We describe Houdini's design in Section 4. We describe the implementation of `Houdini` in Section 5. We briefly describe the container attack surface analysis in Section 6. In Section 7 we present case studies showing how `Houdini` can detect basic misconfigurations, and explain the tricks `Houdini` currently implements and their associated vulnerabilities. Section 9 concludes our paper.

2 Motivation

2.1 The Container Confinement Problem

Container confinement can be defined as the set of mechanisms and strategies used to isolate a containerized application from the host system and other containers. Effective confinement is critical for maintaining system security and stability, as it ensures that containers operate within defined boundaries by restricting their access to system resources, processes, networks, inter-process communication (IPC), filesystems, devices, and kernel interfaces.

A challenge with container confinement is that the level of confinement of a container is not fixed; it can vary based on

the specific requirements or intentions of the user. In other words, the degree to which a container is isolated from the host system or other containers is customizable to meet the user's needs. At one end of the spectrum, you can lock a container down so tightly that it is almost entirely isolated from the host system, which minimizes potential attack surfaces but can also render the container incapable of performing key tasks. On the other hand, allowing a container too much freedom undermines the security benefits of containerization by exposing more privileges than necessary. The optimal configuration, however, is one that strikes a balance between maximum confinement (restricting access to system resources as much as possible) and functional flexibility (enabling the container to perform its intended tasks). Achieving this balance is the core problem and goal of container confinement.

2.2 The Semantic Gap Issue

There exists a semantic gap in containerized systems, where the intended security policies and the actual behavior of the system don't always align. The semantic gap refers to the difference between what a system is meant to do and how it actually behaves in practice. In other words, it's the disconnect between the theoretical design of the system and its real-world execution. This gap can occur for various reasons, such as miscommunications, misunderstandings, or limitations in the system's design or configuration.

For example, when setting up a security system, an administrator may want to ensure that containers in a cloud environment are completely isolated from each other and the host system. The goal is to set up strict access controls so that containers cannot interact with each other or the host. However, due to complex settings, misconfigurations, or flaws in how security mechanisms (like namespaces, cgroups, SELinux, etc.) interact, the system might not enforce these isolation policies correctly. As a result, the actual performance may not align with the intended security model.

This gap can happen in many systems, but it is especially common in complex environments like containers, where multiple interdependent components (e.g., the container runtime, kernel features, and security tools) must work together. Misunderstandings or mistakes in configuring these components can lead to unexpected issues or security vulnerabilities, even when the system seems correctly set up. Another factor contributing to the semantic gap is the interaction between different security layers within a containerized system. Security mechanisms such as SELinux, AppArmor, seccomp, cgroups, and network policies must work together without conflicting or creating unintended loopholes. However, in practice, compatibility issues can arise. For instance, some applications require specific capabilities to function, leading administrators to disable security restrictions to get them working, which introduces unintended security risks. Ensuring that different security mechanisms complement rather than contradict each other is essential to closing the semantic gap.

3 Related Work

At the time of this writing, to our knowledge, `Houdini` is the first comprehensive approach for evaluating and comparing *container confinement mechanisms*. Much of the container security literature to date has focused on either vulnerability scanning of container images, building offensive/attack tools for container escapes, or proposing best practices for securing container runtimes. While these papers, tools, and documents do not directly share our research objectives, they still broadly fit into the container security landscape. Additionally, research on VM escapes and sandbox escapes has provided valuable insights into isolation vulnerabilities, many of which relate to container security. Techniques used to break out of hypervisors and sandboxed environments often share similarities with container escape methods, which the importance of strong confinement mechanisms. The remainder of this section does not aim to exhaustively enumerate all container security tools and systems, but rather to highlight the various salient methods, including prior research on times the confinement of VMs and sandboxes have been escaped as it is related to work.

Vulnerability Scanners. Many free, open, and closed-source tools exist for identifying the presence of known vulnerabilities in container images. Shu *et al.* [6] developed a vulnerability scanning tool to scan Docker Hub container images at scale. They found that on average, official and community images have concerning amounts of vulnerabilities (180+), and that these vulnerabilities remain unpatched for hundreds of days. Their study utilized the Docker Image Vulnerability Analysis (DIVA) framework, which automates the discovery, download, and analysis of over 300,000 Docker images. The findings revealed that more than 80% of both official and community images contained at least one high-severity vulnerability, and many images were not updated for extended periods. Moreover, vulnerabilities were often inherited from parent images to child images, further increasing security risks.

Best Practices. Users in search of (often high-level) advice on how to improve the security of their container deployments can consult one of many best practices guides available online. One such guide is the Center for Interent Security (CIS) Docker Benchmark [1]. In our experience, guides tend to offer generic advice (e.g., "*Only allow read access to the root filesystem*". This is problematic because it oversimplifies container security. Generic advice like "Only allow read access to the root filesystem" may seem straightforward, but it doesn't account for the specific context or unique configuration of an individual deployment. Such one-size-fits-all recommendations can lead users to believe that simply following these broad guidelines is enough to secure their systems, potentially leaving critical vulnerabilities unaddressed. In reality, effec-

tive container security requires a nuanced, tailored approach that considers the particular needs and threat models of each environment. Houdini fills the gap between generic advice and practical security by enabling you to objectively evaluate the effectiveness of your container confinement measures.

Offensive Tools. The offensive security community has developed a range of specialized tools aimed at identifying and exploiting vulnerabilities within containerized environments. Containers, while providing strong isolation, are not immune to attacks, and these tools focus on testing the robustness of container defenses. CDK[1] and DEEPCE[2] are two widely used, open-source container penetration testing toolkits that leverage a collection of known exploits to gain persistence, escape the container, and gather sensitive information about the container environment. These tools aim to bypass isolation mechanisms like namespaces, cgroups, and SELinux, often using a variety of aggressive techniques to identify weaknesses in container security. CDK specifically uses any available method to circumvent security measures, providing a broad toolkit to test the limits of a container's defense. Meanwhile, DEEPCE, developed by stealthcopter, focuses on privilege escalation and container enumeration, attempting to exploit specific weaknesses to achieve container escapes, thereby compromising the host system. In contrast, Houdini takes a more systematic and structured approach to container security testing. Rather than focusing solely on exploiting vulnerabilities, Houdini allows for the direct comparison of defensive techniques by clearly defining the testing environment, outlining the exploit steps, and documenting the expected outputs.

Virtual Machine Escapes VM escapes occur when an attacker successfully breaks out of a virtualized confined environment, gaining unauthorized access to the host system. Given that hypervisors manage multiple VMs on shared hardware, such escapes can lead to host compromise, cross-VM attacks, and lateral movement within cloud environments. One of the most notable VM escape vulnerabilities is VENOM (CVE-2015-3456) [7], which exploited a flaw in the virtual floppy disk controller (FDC) implementation in QEMU, affecting platforms like Xen, KVM, and VirtualBox. Other significant VM escape vulnerabilities include Cloudburst (2009), which exploited VMware's virtual graphics driver; CVE-2018-6981 and CVE-2018-6982, which targeted VMware ESXi, Workstation, and Fusion [8]; and multiple VirtualBox vulnerabilities (2017) enabling guest-to-host code execution [9]. Additionally, side-channel attacks like MDS/ZombieLoad (2019) demonstrated how speculative execution flaws could be leveraged across VMs [10]. More recent research has uncovered weaknesses in encrypted virtualization, such as the SEVered attack (2018), which extracted plaintext memory from AMD SEV-protected VMs [11], and the Heckler attack (2024), which manipulated non-timer interrupts to compromise confidential VMs in AMD SEV-SNP and Intel TDX [12].

Sandbox Escapes Sandboxing is a critical security mechanism for restricting the execution of untrusted code, commonly used in browsers, mobile apps, and security tools. A sandbox escape occurs when an attacker bypasses these restrictions, executing arbitrary code with higher privileges and potentially leading to full system compromise. Modern web browsers heavily rely on sandboxing to isolate JavaScript execution from the underlying OS. However, several high-profile vulnerabilities have enabled sandbox escapes. For example, CVE-2019-5786 [13], a use-after-free vulnerability in the FileReader API of Google Chrome, allowed attackers to escape Chrome's sandbox and execute code on the host. Other notable sandbox escape vulnerabilities include CVE-2022-26706, which exploited a flaw in macOS's App Sandbox to execute code outside of restricted permissions [14]; CVE-2023-36719, a stack corruption vulnerability in Windows OS that enabled escape from the Chromium sandbox [15]; and CVE-2024-4040, which was actively exploited to break out of CrushFTP's Virtual File System (VFS) [16]. Additionally, vulnerabilities in Judge0, an online code execution system, exposed systems to full takeover by allowing attackers to escape its sandbox [17]. Research has also uncovered critical flaws in vm2, a JavaScript sandbox library, which could be exploited to bypass its security restrictions and gain remote code execution [18].

4 Design

Houdini is built using Python 3 and leverages a QEMU KVM VM to provide a controlled testing environment for container security validation. Instead of running containers directly on the host system, Houdini spins up a dedicated VM, ensuring that any security tests remain isolated and do not affect the host system.

At the core of Houdini's architecture, a Flask-based server runs on the host operating system. This server acts as a bridge between the host and the QEMU KVM VM, facilitating communication between them. The Flask server is responsible for sending commands and configuration data to the guest VM and receiving trick execution results. When a trick is initiated, the Houdini client on the host instructs the VM to start a new trick. The VM then launches a container with the specified security configurations and executes the trick inside it.

Once the trick completes, the VM collects the results and sends them back to the host via the Flask server. The host then processes this data, logs the results, and determines whether the trick passed. A trick in `Houdini` is a structured security check designed to evaluate whether a containerized environment enforces confinement properly. Each test systematically probes specific security controls (e.g., namespaces, cgroups, capabilities) to determine if a container remains within its intended isolation boundaries or if a misconfiguration, vul-

[1] https://github.com/cdk-team/CDK
[2] https://github.com/stealthcopter/deepce

nerability, or unexpected behavior allows for unintended behaviour such as, but not limited to privilege escalation, resource leakage, or escape.

This structured workflow ensures that Houdini remains independent of the host system's container runtime while providing a reliable and repeatable environment for testing container confinement mechanisms.

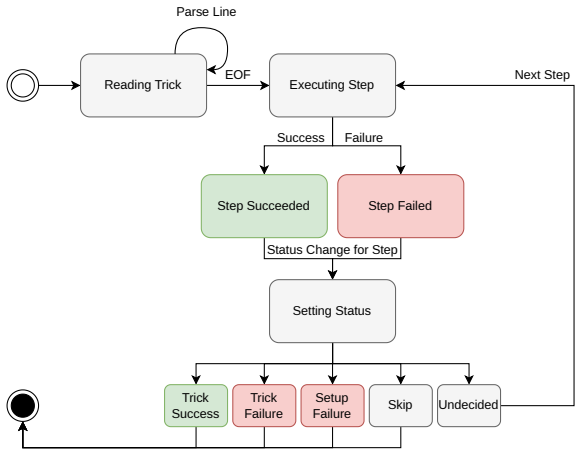


Figure 1: A state machine diagram of Houdini running a Trick.

4.1 Design Goals

`Houdini` is designed to provide a systematic and reproducible approach to testing container confinement mechanisms. One of its primary goals is to evaluate whether container security features, such as namespaces, cgroups, seccomp, and Linux capabilities are effectively enforcing isolation. Many security tools rely on theoretical guarantees or static analysis, but `Houdini` focuses on empirical validation by executing controlled security tests that verify the correct enforcement of security mechanisms. By doing so, `Houdini` ensures that container security mechanisms are correctly configured.

One of the goals of `Houdini` is not only for testing container isolation but also as a means to optimize container configurations. By systematically simulating various tasks within Docker containers, Houdini can help determine the minimal set of privileges required for a container to perform its intended function. For example, administrators can start with a container configured with broad privileges and then incrementally reduce permissions, such as Linux capabilities, seccomp filters, or resource limits, while monitoring task success. Houdini's testing framework will identify the point at which functionality is maintained while unnecessary privileges are removed, thereby pinpointing the least privilege configuration that still allows the container to operate effectively. This approach is highly beneficial because it adheres to the principle of least privilege. Minimizing the privileges granted to a container reduces its attack surface, making it significantly harder for an attacker to exploit vulnerabilities.

Although runtime monitoring and anomaly detection, such as with eBPF-based tracing can provide immediate insights into container behavior, our goal isn't to implement an exhaustive, continuous surveillance system. Instead, we aim to focus on auditing and validating the underlying isolation mechanisms (like cgroups, namespaces, and mountpoints) to ensure they are correctly configured and consistently enforced.

The following are additional goals of Houdini:

Verifying Whether Security Mechanisms Are Truly Working as Expected: Many security mechanisms in containerized environments, such as Linux namespaces, cgroups, seccomp filters, and capabilities are intended to enforce process isolation and prevent privilege escalation. However, just because a security feature is enabled does not necessarily mean it is functioning correctly. Misconfigurations, improper runtime enforcement, and subtle implementation flaws can lead to unexpected security gaps. `Houdini` helps verify whether these security mechanisms are actually enforcing the intended restrictions by executing controlled tests (tricks) that attempt to bypass or manipulate confinement policies. If a test succeeds in breaking out of isolation, it provides direct evidence that the security mechanism is not functioning as expected.

Identifying Gaps Between Intended Confinement Policies and Actual Enforcement: There is often a discrepancy between security policies that administrators configure and what is actually enforced at runtime. For example, a security policy may specify that a container should not have network access, but due to a misconfigured Docker setting, the container may still be able to establish outbound connections. Similarly, a container may be configured with resource constraints (e.g., CPU and memory limits via cgroups), but in practice, those constraints may be ineffective or bypassable. `Houdini` systematically tests these security policies against real execution environments to uncover where enforcement fails. By detecting these discrepancies, organizations can adjust their configurations and improve security postures, ensuring that containers remain properly confined under expected conditions.

Challenging Misleading Security Assumptions with Empirical Testing: Container security is often discussed in terms of theoretical guarantees, with many security best practices based on assumptions rather than direct testing. For instance, documentation may state that a container running in unprivileged mode cannot escalate to root, but Houdini actively tests whether that is actually the case under different container configurations. By running real-world security validation instead of relying solely on security claims, `Houdini` helps challenge misleading assumptions and exposes where security mechanisms fail in practice. This is particularly important as

container escapes and privilege escalation exploits continue to emerge, demonstrating that security controls do not always work as expected. By grounding security claims in measurable results rather than theory, Houdini helps researchers and practitioners develop more robust container security strategies.

In order to achieve the aforementioned goals, we designed Houdini with the following in mind:

1. *Reproducible Results.* A key goal of Houdini is to ensure that tests yield consistent and repeatable results. By controlling system variables and maintaining a structured testing environment, Houdini allows researchers and practitioners to compare results across different system configurations and security policies.

2. *Separation from the Host.* Houdini is designed to run tests inside a controlled containerized environment within a virtual machine (VM), ensuring that even if a test triggers a security vulnerability, it does not compromise the host system running the tests. This design choice enhances the safety of security evaluations, preventing unintended side effects on the underlying infrastructure.

3. *Test Case Expressiveness.* Houdini test cases (called "tricks") should be maximally expressive, such that some combination of steps can be used to achieve and test any desired result. It should be possible to define a new Houdini trick and modify existing tricks without modifying the Houdini binary. Moreover, it should be clear from reading a defined trick precisely what steps are involved, the consequences of each step passing or failing, and the overall nature of the exploit being tested.

4. *Focus on Observing Failures, Not Preventing Them.* Houdini is built to identify security weaknesses rather than defend against attacks.

4.2 Security of the Testing Environment

If Houdini were compromised during testing, it would not necessarily pose a threat to the host system due to the design of the testing environment. Houdini is intentionally designed to operate within a controlled, containerized environment inside a VM. This means that even if a test were to exploit a vulnerability within the containerized environment, the impact would be contained within the VM, isolated from the host system and other containers. The VM itself serves as an additional layer of security, providing a buffer that prevents the compromise from reaching the host infrastructure.

5 Implementation

Exploits that Houdini can test are divided into three separate files: a configuration file, a Python file, and a Dockerfile. Each Houdini trick begins with a configuration file (see Listing 1), which contains the Docker container's settings and environment configuration. The configuration file defines various Docker-specific parameters that are directly fed to the Docker API to start a container with the desired settings. Once the configuration is parsed, Houdini communicates these settings to the Docker API, using Docker's docker run command to initiate the container with these parameters. This ensures that the container is launched with the exact environment specified. This automation ensures the container is started exactly as specified, and that parameters such as network mode, mounts, security options, and user privileges are all controlled from a single file. This makes it easier to test, replicate, and share container environments.

Listing 1: Configuration file for a default Docker out-of-the-box security settings applied when a container is launched.

```
name: out-of-the-box docker configuration
dockerfile:
  - path: Dockerfile
dependencies:
  - server: False
docker_config:
  - network_mode: bridge
  - read_only: False
  - security_opt: []
  - pid_mode: null
  - cpu_shares: 1024
  - volumes : []
  - mem_limit: null
  - cpuset_cpus: null
  - cpu_quota: null
  - cpu_period: null
  - cap_add: []
  - cap_drop: []
  - privileged: False
  - user: root
  - pids_limit: null
  - ipc_mode: private
trick:
  - path: /houdini/tricks/HostMount
```

The second component is a Python file that contains the trick implementation. This Python file is executed inside the container and is responsible for carrying out the trick scenario.

Finally, the Dockerfile is used to configure and set up the container's environment. It specifies how the container should be built, including copying necessary files into the container and installing required dependencies. The Dockerfile also configures the container with the settings defined in the configuration file, ensuring the container is set up correctly to run the trick.

Listing 2: Exploit for CVE-2024-21616

```
import os

relative_path = '../../../../../'

try:
    os.chdir(relative_path)
    print(f"Successfully changed directory to: {os.getcwd()}")
except FileNotFoundError as e:
    print(f"Error: {e}")
except PermissionError as e:
    print(f"Error: {e}")
except Exception as e:
    print(f"Unexpected error: {e}")
```

Listing 3

```
FROM ubuntu:20.04
RUN apt-get update -y
WORKDIR /proc/self/fd/9
CMD ["bash", "-c", "ls
↪ ../../../../../houdini/tricks/HostMount"]
```

This modular design, comprising of the configuration file, Python script, and Dockerfile provides flexibility to define and test various container escapes and vulnerabilities. It also allows users to easily customize or extend existing tricks by modifying any of the three components to suit their needs.

6 Container Attack Surface Analysis

In containerized environments, security challenges often arise from vulnerabilities within specific components of the architecture. As such, understanding the relationships between the different components, such as container registries, images, runtime, namespaces, cgroups, and various security modules like eBPF, seccomp, and Linux capabilities is crucial for identifying potential attack surfaces. This is where a comprehensive architecture diagram comes into play.

Figure 2 outlines the key components of a containerized system, highlighting the critical paths and interactions that security tests should focus on. Security testing should aim to cover these components extensively to ensure that vulnerabilities in any of these layers are adequately addressed. However, some components, like network interfaces, system calls, and cgroups, are more frequently targeted in real-world attacks and thus warrant more attention in testing. The diagram will guide our understanding of which components to prioritize during testing, ensuring that our coverage is both thorough and effective.

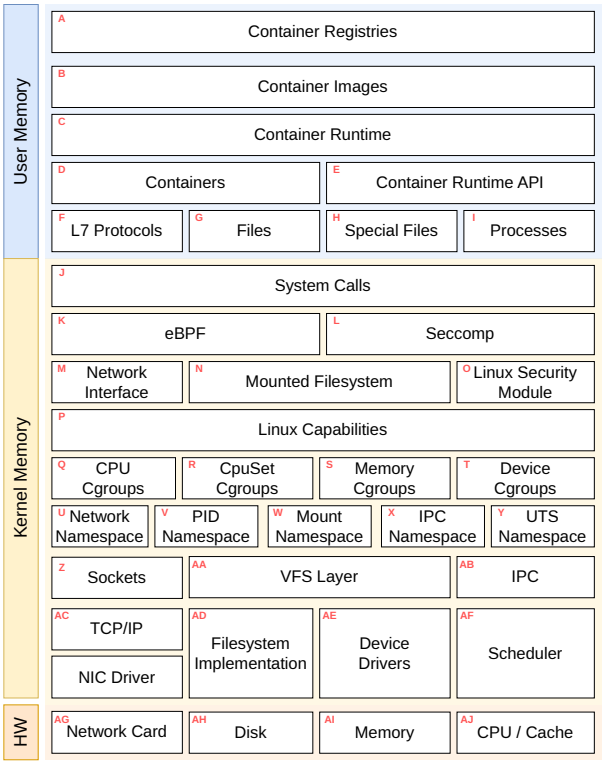


Figure 2: An architectural diagram of a container deployment environment depicting the attack surface created by its various components.

7 Evaluation

To assess the effectiveness of container confinement mechanisms, we conducted a series of controlled experiments to evaluate whether common security policies successfully prevent resource starvation, unauthorized access attempts, and container escapes. Our evaluation focuses on testing various isolation features, including namespaces, cgroups, seccomp, and Linux capabilities, under different configurations and attack scenarios. These are done in the form of three case studies.

7.1 Case Study 1 - Evaluating Host Process Manipulation via PID Namespace Escape

7.1.1 Introduction to the Evaluation

In this evaluation, we investigate whether a Docker container running with `-pid=host` can interact with and manipulate host processes. The PID namespace is responsible for process isolation, ensuring that containerized processes remain separate from host processes. However, when `-pid=host` is enabled, the container shares the host's PID namespace, making all host processes visible inside the container. We

assess the security implications of this setting by evaluating the following:

1. Whether a container with `-pid=host` can list host processes.
2. Whether the container can send signals `SIGKILL`, `SIGTERM` to host processes.
3. Whether enabling `CAP_SYS_ADMIN` allows process tracing of host processes.
4. Whether a container can read host process memory via `/proc/$PID/mem`.
5. Whether a container with `-pid=host` can manipulate processes running inside another container that also has `-pid=host`.

7.1.2 Experimental Setup

The experiments were conducted using Docker running inside a QEMU VM to ensure an isolated and controlled test environment. The system specifications were as follows:

- **Host Machine:** 20 Core CPU, 16GB RAM, Ubuntu 24.04
- **Virtual Machine:** QEMU, 4 Core CPU, 12GB RAM, No Distibution
- **Docker Version:** 27.0.3
- **Kernel Version:** Linux 6.1
- **Container Runtime:** `runc` 1.1.11

The container was started with various security configurations:

- Default Docker container configuration (separate PID namespace)
- Docker container with `-pid=host`
- Docker container with `-pid=host` and additional capabilities `CAP_KILL`, `CAP_SYS_PTRACE`
- Privileged Docker container with `-pid=host`
- Two separate containers, both with `-pid=host`, to evaluate cross-container process manipulation.

For each security configuration, different methods were tested to determine whether a process inside the container could interact with host processes or another container's processes.

7.1.3 Security Testing

The evaluation began by assessing whether a container with `-pid=host` could list host processes using `ps aux`. As expected, the container displayed all host processes, confirming that sharing the host PID namespace fully exposes the host process table. Following this, we tested whether the container could send termination signals (`SIGKILL`, `SIGTERM`) to host processes using `kill -9 <PID>`. Without the `CAP_KILL` capability, the operation was not permitted. However, enabling `CAP_KILL` allowed process termination within the namespace's limits, and granting `-privileged` access enabled the container to kill any host process. Next, we investigated whether the container could attach a debugger to host processes via `strace -p <PID>`. Without `CAP_SYS_PTRACE`, the operation was denied, but enabling this capability allowed successful attachment to a host process, raising concerns about potential leakage of sensitive data through system call tracing. We then examined whether the container could read host process memory by attempting to access `/proc/<host_pid>/mem`. With Docker's default restrictions, the operation was denied, but when running the container with `-privileged`, the process memory could be read using `cat /proc/<host_pid>/mem`. Finally, we tested cross-container process manipulation by launching two containers with `-pid=host`. A process running in Container A (e.g., `sleep 1000`) was visible to Container B, which was able to list and terminate it, proving that containers sharing the host PID namespace are not isolated from each other and can manipulate each other's processes. This finding underscores the risk of using `-pid=host`, as it allows one compromised container to interfere with others running on the same host.

The result was that the process inside Container A was visible to Container B, and it could be terminated using `kill -9 <PID>`, proving that containers with `-pid=host` share the same namespace and are not isolated from each other.

7.1.4 Results and Findings

To summarize, using `ps aux`, containers with `-pid=host` could list all host processes. Sending signals `SIGKILL`', 'SIGTERM was blocked without `CAP_KILL`, but enabling the capability allowed termination of processes within the namespace's limits. Debugging `strace -p <PID>` required `CAP_SYS_PTRACE`, and reading host memory `cat /proc/<PID>/mem` was blocked unless the container was privileged. Finally, two containers with `-pid=host` were able to list and terminate each other's processes, demonstrating cross-container process interference. The effectiveness of different configurations in preventing host and cross-container process manipulation is also illutrated in Table 1.

Table 1: Effectiveness of Docker Security Mechanisms Against Host Process Manipulation

Security Configuration	List Host Processes	Send Signals	Attach Debugger	Read Memory	Cross-Container Manipulation
Default Docker (PID namespace isolation)	✗	✗	✗	✗	✗
`-pid=host` (no extra caps)	✓	✗	✗	✗	✓
`-pid=host` + `CAP_KILL`	✓	✓	✗	✗	✓
`-pid=host` + `CAP_SYS_PTRACE`	✓	✗	✓	✗	✓
Privileged Container (`-pid=host` + `-privileged`)	✓	✓	✓	✓	✓

7.1.5 Key Takeaways

- `-pid=host` completely breaks process isolation, allowing full visibility into host and cross-container processes.
- Containers sharing the host PID namespace can manipulate each other's processes.
- Privileged containers with `-pid=host` have unrestricted access to host process memory, debugging, and termination.

7.2 Case Study 2 - Evaluating Docker Port Forwarding Restrictions

7.2.1 Introduction to the Evaluation

In this case study, we evaluate Docker's network isolation mechanisms to determine how well they enforce port binding restrictions. Specifically, we examine whether privileged ports (1-1024) can be bound by containerized processes, and how changes in Docker's security model (pre- and post-20.03) have affected these restrictions.

Starting with Docker version 20.03, a significant change was introduced. Docker redefined unprivileged ports to start at 0 instead of 1024, effectively removing the need for `NET_BIND_SERVICE` to bind to privileged ports. This has direct security implications, as it potentially allows attackers with root access inside a container to bind services to privileged ports that should traditionally require special permissions.

Thus, we evaluate two key aspects of Docker's security model:

- **Docker $\leq$ 20.03:** To bind to privileged ports, root users are to use required NET_BIND_SERVICE.
- **Docker $\geq$ 20.03:** Root users can bind to privileged ports without 'NET_BIND_SERVICE'.

7.2.2 Experimental Setup

The experiments were conducted in a controlled QEMU-based virtualized environment, ensuring reproducibility and isolation from external interference.

- **Host Machine:** 20 Core CPU, 16GB RAM, Ubuntu 24.04
- **Docker Versions Tested:**
  - 19.03.14 (pre-20.03)
  - 24.0.2 (post-20.03)
- **Container Runtime:** runc 1.1.7
- **Containers Used:**
  - Default container (bridge network mode, no additional privileges)
  - Container with 'NET_BIND_SERVICE'
  - Custom seccomp profile restricting 'bind' syscall
  - Privileged mode

Each test involved binding an HTTP server to port 23 (a privileged port), under different security configurations.

7.2.3 Security Testing

1. Port Binding Restrictions in Docker < 20.03
Test Procedure:

- A containerized HTTP server was started, attempting to bind to port 23.
- Tests were conducted under varying privilege levels:
  - With and without 'NET_BIND_SERVICE'
  - As root vs. non-root
  - With a seccomp profile denying the 'bind' syscall

Table 2: Effectiveness of Docker Security Mechanisms in Port Binding (Docker < 20.03)

Security Configuration	Result
`NET_BIND_SERVICE + root`	✓
`root`	✗
`NET_BIND_SERVICE + custom Seccomp (deny bind syscall) + root`	✗
`NET_BIND_SERVICE + non-root`	✗

2. Port Binding Restrictions in Docker $\geq$ 20.03
Test Procedure:

- The same HTTP server was used, attempting to bind to port 23.
- We evaluated how removing 'NET_BIND_SERVICE' affected the ability to bind to privileged ports.

Table 3: Effectiveness of Docker Security Mechanisms in Port Binding (Docker $\geq$ 20.03)

Security Configuration	Result
`root`	✓
`Custom Seccomp (deny bind syscall) + root`	✗
`non-root`	✗

7.2.4 Key Takeaways

- **Docker < 20.03:** 'NET_BIND_SERVICE' was required to bind to privileged ports.
- **Docker $\geq$ 20.03:** Root containers can bind to privileged ports without 'NET_BIND_SERVICE'.

8 Case Study 3: Evaluating CVE-2024-21616

8.1 Introduction

CVE-2024-21616 is a critical vulnerability affecting Docker and runc that allows malicious containers to escape their isolation and execute arbitrary commands on the host system. This poses a significant risk to enterprises relying on containerization, as attackers could gain unauthorized access to sensitive data and system resources.

The root cause of this vulnerability is the improper handling of a container's working directory when set to a special file descriptor path such as `/proc/self/fd/<fd>`, where `<fd>` typically points to `/sys/fs/cgroup`. If exploited, this flaw allows a containerized process to break out of its confined environment, effectively compromising the host machine. The vulnerability affects runc versions from `v1.0.0-rc93` to `1.1.11`.

In this case study, we leverage Houdini to systematically assess the effectiveness of Docker's security mechanisms in mitigating CVE-2024-21616. By integrating the vulnerability as a Houdini trick, we analyze the conditions under which the exploit succeeds and evaluate different security configurations.

8.2 Methodology

To test the exploitability of CVE-2024-21616, we designed a Houdini trick that executes the following steps:

1. Launch a container with a vulnerable version of runc.
2. Set the working directory to `/proc/self/fd/<fd>`, pointing to a controlled file descriptor.
3. Attempt to break out of the container by leveraging the misconfigured working directory.
4. Validate whether commands can be executed on the host system.

8.3 Experimental Setup

We conducted tests across three different configurations:

1. **Default Docker Configuration:** No additional security measures.
2. **Hardened Configuration:** Using seccomp, AppArmor, and dropping unnecessary capabilities.
3. **Patched Environment:** Upgrading runc to version `1.1.12` or later.

8.4 Results and Findings

Table 4: Effectiveness of Docker Security Configurations Against CVE-2024-21616

Configuration	Exploit Success
Default Docker (runc 1.1.11)	✓
Hardened Security Profile (runc 1.1.11)	✓
runc 1.1.12+	✗

Our findings reveal that neither Docker's default security settings nor additional confinement mechanisms (e.g., seccomp, AppArmor) were effective in preventing exploitation. The only mitigation strategy that successfully blocked the attack was upgrading runc beyond version `1.1.11`, as the vulnerability lies within the container runtime itself.

9 Conclusion

In this paper, we presented Houdini, a security benchmarking framework designed to evaluate the effectiveness of container confinement mechanisms. As container-based workloads continue to dominate cloud infrastructures, ensuring that containers are properly isolated is paramount for maintaining security. Houdini offers a systematic and empirical approach to test and assess various container security configurations, highlighting potential misconfigurations that may compromise container isolation. Our evaluation demonstrated the ability of Houdini to detect critical misconfigurations, including those leading to privilege escalation. By providing a reliable and reproducible method for testing container confinement, Houdini helps bridge the gap between theoretical security assumptions and practical real-world behavior.

As container technologies evolve, it is essential to continuously assess and improve their security mechanisms. Houdini's modular design ensures that it remains adaptable to new security challenges and configurations, allowing for ongoing validation and refinement of container security practices. By offering an open framework for container confinement benchmarking, Houdini aims to support the development of more secure container technologies, ultimately contributing to a safer cloud infrastructure for the future.

References

[1] Cis docker benchmark. In *CIS Docker Benchmark*, 2024.

[2] Sari Sultan, Imtiaz Ahmad, and Tassos Dimitriou. Container security: Issues, challenges, and the road ahead. *IEEE access*, 7:52976–52996, 2019.

[3] John L Henning. Spec cpu2000: Measuring cpu performance in the new millennium. *Computer*, 33(7):28–35, 2000.

[4] Raghunath Nambiar, Nicholas Wakou, Forrest Carman, and Michael Majdalany. Transaction processing performance council (tpc): State of the council 2010. In Raghunath Nambiar and Meikel Poess, editors, *Performance Evaluation, Measurement and Characterization of Complex Systems*, pages 1–9, Berlin, Heidelberg, 2011. Springer Berlin Heidelberg.

[5] Martin Zboril and Vlasta Svatá. Performance comparison of cloud virtual machines. *Journal of Systems and Information Technology*, 2025.

[6] Rui Shu, Xiaohui Gu, and William Enck. A study of security vulnerabilities on docker hub. In *Proceedings of the Seventh ACM on Conference on Data and Application Security and Privacy*, pages 269–280, 2017.

[7] Cheick Abdoul-Kader and Shih-Hao Chang. A novel venom attack identification mechanism in cloud virtualization environment. *Communications of the CCISA*, 24(1):61–72, 2018.

[8] VMware Security Team. Vmware virtual machine escape vulnerabilities: Cve-2018-6981 and cve-2018-6982. 2018. Accessed on: March 3, 2025.

[9] Google Project Zero. 10 new vm escape vulnerabilities discovered in virtualbox. 2017. Accessed on: March 3, 2025.

[10] Stefan Lipp, Michael Schwarz, and Daniel Gruss et al. Mds: Microarchitectural data sampling. *IEEE Symposium on Security and Privacy*, 2019. Accessed on: March 3, 2025.

[11] Felix Dörre and Florian Piessens. Severed: Subverting amd's virtual machine encryption. *27th USENIX Security Symposium*, 2018. Accessed on: March 3, 2025.

[12] Amit M. and Thomas Eisenbarth. Heckler attack: Breaking amd sev-snp and intel tdx with non-timer interrupts. *arXiv preprint arXiv:2404.03387*, 2024. Accessed on: March 3, 2025.

[13] Philippe Laulheret. Analysis of a chrome zero day: Cve-2019-5786. 2019. Accessed on: March 3, 2025.

[14] Jonathan Bar Or. Uncovering a macos app sandbox escape vulnerability: A deep dive into cve-2022-26706, 2022. Accessed on: March 3, 2025.

[15] James Forshaw. Escaping the sandbox: A bug that speaks for itself, 2023. Accessed on: March 3, 2025.

[16] Satnam Narang. Cve-2024-4040: Crushftp virtual file system (vfs) sandbox escape vulnerability exploited, 2024. Accessed on: March 3, 2025.

[17] Ravie Lakshmanan. Sandbox escape vulnerabilities in judge0 expose systems to complete takeover, 2024. Accessed on: March 3, 2025.

[18] FortiGuard Labs. Vm2 sandbox escape vulnerability | outbreak alert, 2023. Accessed on: March 3, 2025.